

Object-Oriented DBMSs

- Standards-based effort
 - ODMG (Object Data Management Group)
- Have produced several important standards so far
- Object Description Language ODL
 - Specifies the kind of schema-construction operations with a similar purpose to what we use in SQL's CREATE TABLE
- Object Query language
 - Provides an interface that attempts to mimic the flexibility and power of SQL, but in an OO framework

The Grand Idea

- ODMG imagines the following scenario:
 - OO-DBMS vendors implement versions of OO languages like C++ or Java
 - These languages have OQL extensions
 - Therefore a programmer could transfer data from the host language to the database (and back) in a seamless manner
- ODL is meant to define **persistent classes**
 - That is, these are objects stored permanently in the database.
 - ODL classes look like Entity sets with binary relationships along with methods.
 - ODL class definitions are part of the extended OO host language

ODL: overview

- Class declarations have several parts.
 1. Name of the class
 2. Optional key declarations
 3. Element declarations
 - An **element** is either an attribute, a relationship, or a method

```
class <name> {  
    <list of element declarations, separated by semicolons>  
}
```

Declarations: attributes; relationships

- Attributes are (normally) elements with a type that does not involve classes

```
attribute <type> <name>;
```

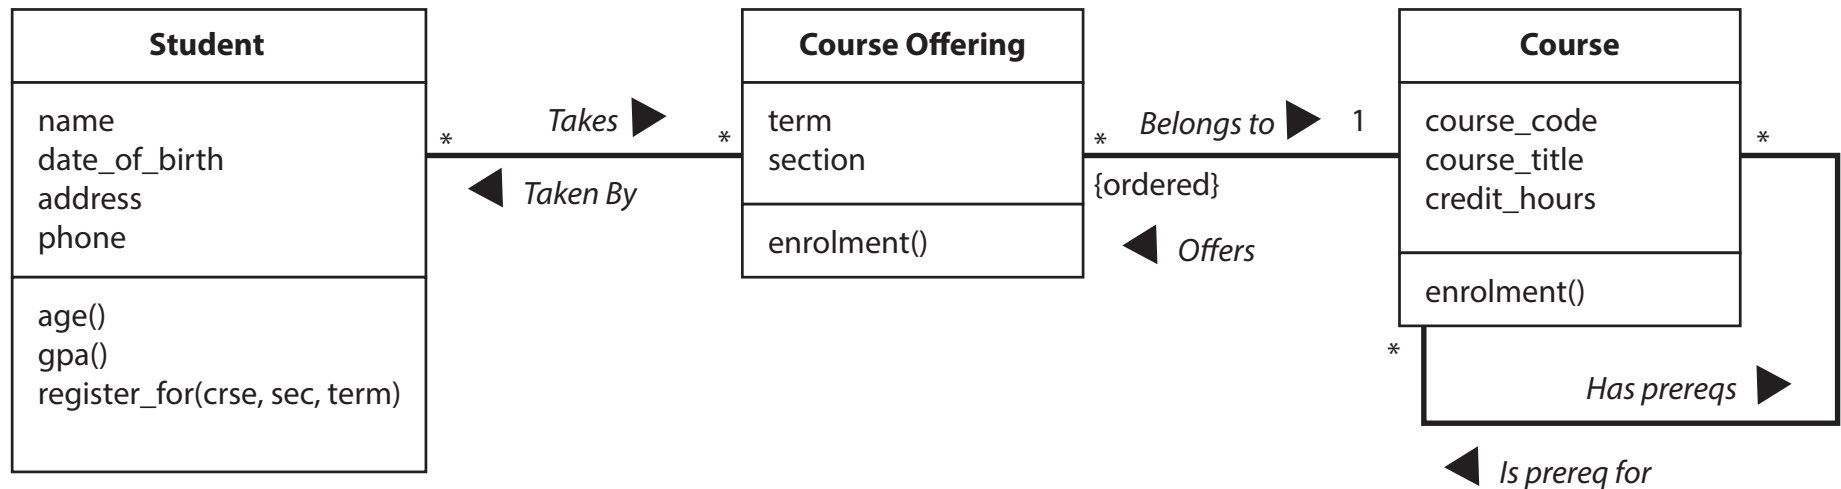
- Relationships connect an object to one or more other objects of one class

```
relationship <type> <name> inverse <relationship>;
```

Inverse relationships

- Given a class **C** that has some relationship **R** to class **D**...
- ... then class **D** must have some relationship **S** to class **C**.
- **R** and **S** must be true inverses.
 - If object **d** is related to object **c** by **R**, then **c** must be related to **d** by **S**.
- ODMG Object model supports only unary and binary relationships

Example: Course system



Example: Pubs and Beers

```
class Pub {  
    attribute string name;  
    attribute string addr;  
    relationship Set<Beer> serves inverse  
}
```

The type of relationship **serves** is a set of Beer objects.

Beer::servedAt;

```
class Beer {  
    attribute string name;  
    attribute string manf;  
    relationship Set<Pub> servedAt inverse Pub::serves;  
}
```

The :: operator connects a name on the right to the context containing that name, on the left.

Relationship phrasing

- There are three possibilities.
 1. A class, like Pub
 - If so, an object with this relationship can be connected to only one Pub object
 2. Set<Pub>
 - The object is connected to a set of Pub objects
 3. Bag<Pub>, List<Pub>, Array<Pub>
 - The object is connected to a bag, list, or array of Pub objects respectively.

Multiplicity of relationships

- Recall that all relationships are either binary or unary
- Many-to-many relationships
 - Have $\text{Set}\langle\ldots\rangle$ for the type of relationship
 - Also for its inverse
- Many-to-one relationships
 - Have $\text{Set}\langle\ldots\rangle$ in the relationship of the **one**...
 - ... and just the class for the relationship of the **many**
- One-to-one relationships
 - Have classes as the type in both directions

Example: Multiplicity

```
class Patron {  
  attribute ...  
  relationship Set<Beer> likes inverse Beer::fans;  
  relationship Beer favBeer inverse Beer::superfans;  
}  
  
class Beer {  
  attribute...  
  relationship Set<Patron> fans inverse Patron::likes;  
  relationship Set<Patron> superfans  
    inverse Patron::favBeer;  
}
```

Many-many uses Set<...>
in both directions.

Many-one uses Set<...>
only with the "one."

Example: Multiplicity

husband and wife are one-one and
inverses of each other.

```
class Patron {  
    attribute ...  
    relationship Patron husband inverse wife;  
    relationship Patron wife inverse husband;  
    relationship Set<Patron> buddies  
        inverse buddies;  
}
```

buddies is many-many and its
own inverse. Note no :: needed
if the inverse is in the same class.

```
class PoliticallyCorrectPatron {  
    attribute ...  
    relationship Patron spouse1 inverse spouse2;  
    relationship Patron spouse2 inverse spouse1;  
    relationship Set<Patron> specialFriends  
        inverse specialFriends;  
}
```

Coping with multiway relationships

- In a word: not very well.
- ODL does not support three-way or higher-degree relationships
- We may simulate multiway relationships
 - Use a **connecting class**
 - Objects represent tuples of objects we would like to connect by the multiway relationship

Connecting classes

- Imagine we want to connect classes **F**, **G** and **H** by some relationship **R**.
- To do this, we create a class **C**:
 - Its object instances represent a triple of objects **(f, g, h)**
 - These are, of course, meant to be instances **F**, **G**, and **H**
- We therefore construct three many-to-one relationships.
 - They go from the **(f, g, h)** object to each of **f**, **g**, and **h**.

Example

- Recall our Pub and Beer classes
 - In an E/R diagram we represented the price at which the beer is sold in the pub.
 - Many-to-many relationships in ODL cannot have a price attribute as is possible in the E/R model.
- One solution:
 - Create class Price...
 - ... and a connecting class PBP ...
 - ... to represent a related pub, beer, and the price.

Example

- Unlike E/R diagrams, attributes in ODL can have a structure
 - Just like C
- We declare this by using some new notation:

```
attribute Struct { <details> } <name of attribute>;
```

```
attribute Enum { <details> } <name of attribute>;
```

- The <details> portion contains:
 - Field name(s) for a Struct
 - Constant(s) for an Enum

Example: Struct and Enum

```
class Pub {  
    attribute string name;  
  
    attribute Struct Addr  
        {string street, string city, string postalCode}  
    attribute Enum Lic  
        {FULL, BEER, NONE}  
    relationships ...  
}
```

Names for the structure
and enumeration

address;

license;

names of the
attributes

Method declarations

- ODL classes may include declarations of methods for the class
 - With this notation, of course, we do specify the method's definition!
- Information for a class:
 - return type (if any)
 - method name
 - argument modes and types (no names!) where modes are **in**, **out** and **inout**
 - any exceptions the method may raise

Example: Methods

```
real gpa(in string) raises (noGrades);
```

1. The method **gpa** returns a **real** number (and we deduce here it is probably some student's GPA)
2. **gpa** accepts a single argument of type **string** (again we deduce it to be the student's name)
 1. The value for the argument will only ever be passed **in**
 2. In a full OO programming language, we would mark the parameter as constant or final
3. **gpa** may raise the exception noGrades

The ODL type systems

- Basic types:
 - int, real/float, string, enumerated types, classes
- Type constructors:
 - Struct for structures
 - Collection types such as Set, Bag, List, Array and Dictionary
- Relationship types can only be either:
 - a class
 - a single collection type applied to a class

Subclassing in ODL

- This is the expected object-oriented notion of subclasses
- We indicate superclass with a colon (':') and the superclass name
 - This notation is borrowed from C++
- A subclass will list only the properties unique to it
 - That is, the subclass will inherit the properties of its superclass (no surprise there)

Example

```
class Ale:Beer {  
    attribute string colour;  
    attribute int    IBU;  
}
```

Keys in ODL

- Unlike E/R you can declare any number of keys for a class
- These keys appear after the class name

```
(key <list of keys> )
```

- A key consisting of more than one attribute needs additional parentheses around the attribute set

Example: Keys

- **name** is the key for Beers

```
class Beer (key name) { ... }
```

- For a Course class:
 - **dept** and **number** form one key
 - **room** and **hours** form another key

```
class Course  
  (key (dept, number),  
   (key (room, hours))) { ... }
```

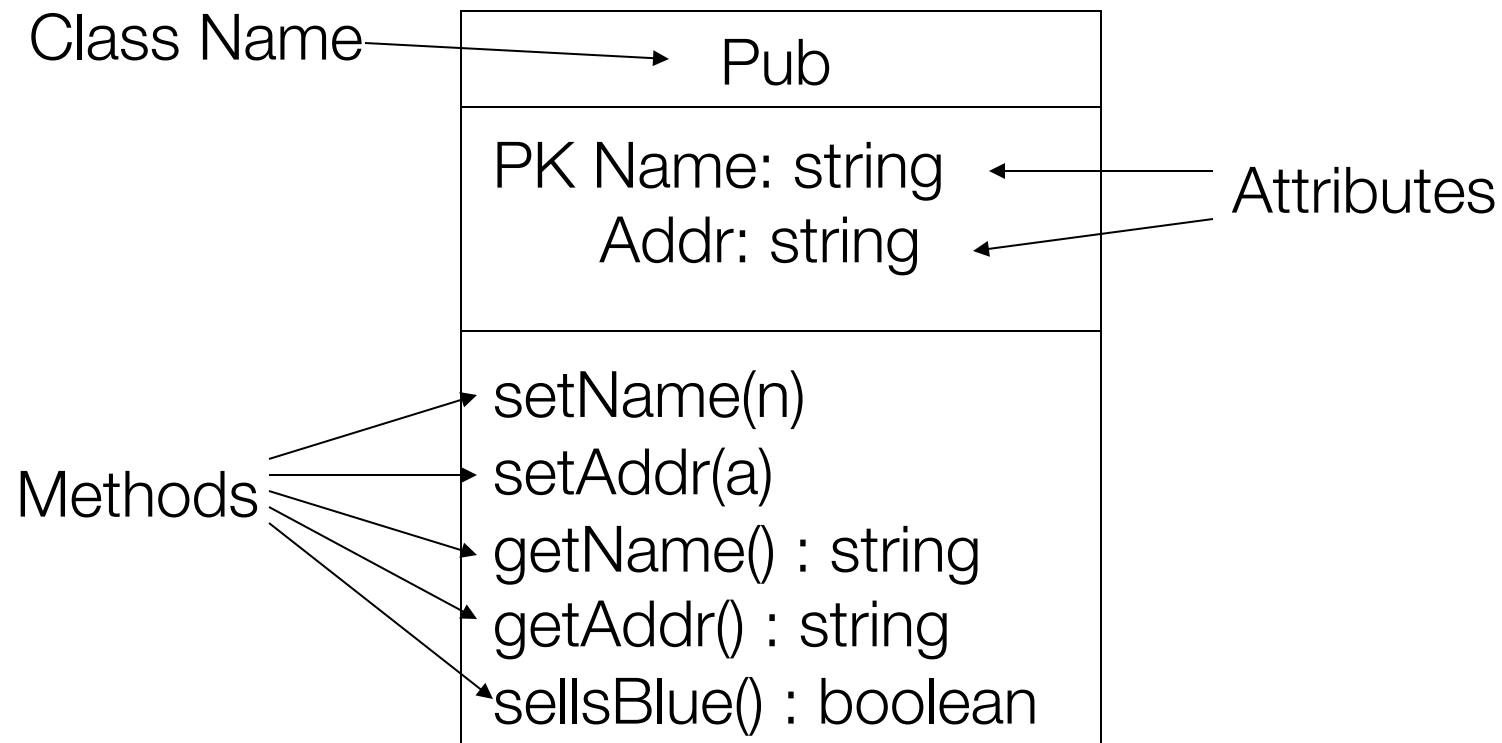
How does UML really fit into all this?

- UML was originally intended for modelling software structure
 - However, it has been adapted as a database modelling notation
- In expressive power, UML lies midway between E/R and ODL
 - No multiway relationships as in E/R
 - Does allow attributes on binary relationships (which ODL does not)
 - Is a graphical notation (unlike ODL)

Classes

- Represent sets of objects
 - Having attributes (i.e., state)
 - Having methods (i.e., behavior)
- Attributes have types
- The notation PK indicates an attributes in the primary key of the object
 - A primary key is optional
- Methods have declarations
 - arguments (if any)
 - return type

Example: Bar Class



Associations




- In UML we already have binary relationships between classes / objects.
- These are represented by named lines
 - There are no diamonds as in E/R
- Multiplicity of the relationship is indicated on either end of the line
 - **m .. n** means from **m** to **n** of these objects **associate with the object on the other end of the line**
 - ***** means **0 to many**
 - **1..*** means **at least one**

Example: Association






Comparison With E/R Multiplicities

E/R




UML




Association Classes

- Attributes on associations are permitted
 - These are **association classes**
 - Similar to attributes on E/R relationships



Subclasses

- Somewhat similar to ISA hierarchies in E/R
 - However, subclass points to superclass with a line ending in a triangle
- Subclasses of a class can be:
 - **Complete** or **partial**
 - **Disjoint** or **overlapping**



Conversion to relations

- We can use any of the three strategies outlined for E/R for conversion
 - Recall the issue is how we translate an entity and its derived entities as relations
- 1. E/R style: each subclass' relation stores only its own attribute, plus key.
- 2. OO-style: relations store attributes of subclass and all superclasses
- 3. Nulls: One relation, with NULLs as needed

Aggregations

- Relationships in UML have an implication
 - Objects on one side may be owned by (or are part of) objects on the other
- This is represented using the diamond notation
 - Diamond appears at the owner's side of the association.
- Further implication here:
 - In a relational schema, owned objects are part of owner tuples.



Compositions

- Like aggregations, but with a further implication
 - Every object must be owned by some object on the other side
 - In OO terms, also means something about when certain objects come into existence and when they cease to exist
- Solid diamond notation
- Often used for subobjects or structured attributes

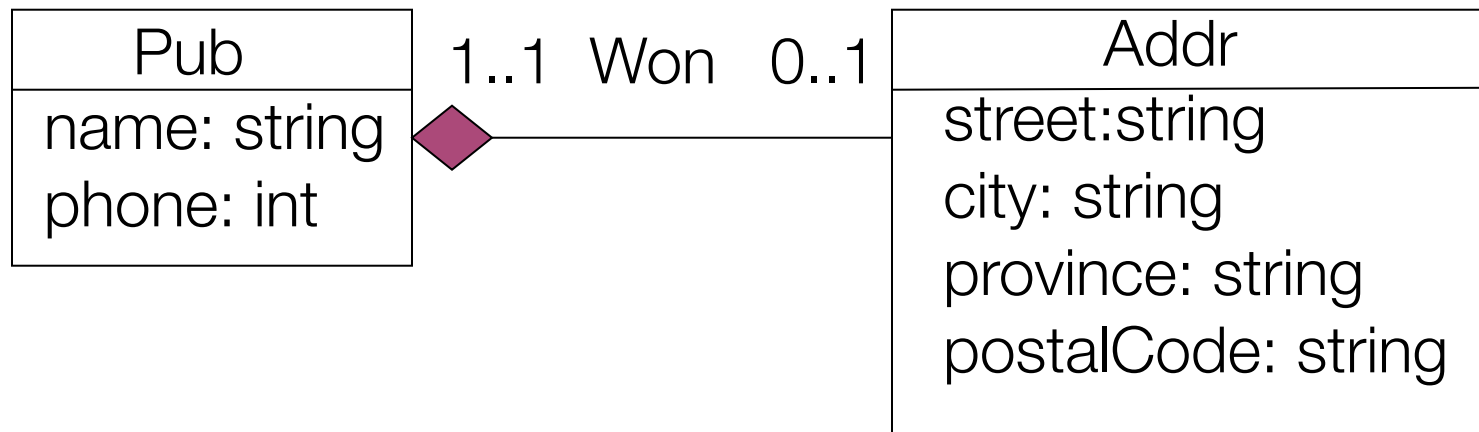


Conversion to relations

- We could store the awards of a beer with the beer tuple itself
- However:
 - This would require an object-relational model (or nested-relation model) for tables.
 - Issue here is no limit exists to the number of awards a beer can win.

Conversion to relations

- Since pub has at most one address...
 - ... it is feasible to add the street, city, province and postal code of Addr to the Pubs relation
- In an object-relational database, Addr can be one attribute of Pubs, with structure



Summary

- E/R model with entities and relationships
- Ternary and higher-order relationships
- Multiplicities
- ISA hierarchies
- Weak entities
- Conversion to relations
- Object Description Language
- Modelling data using UML (and distinctions from ER)

Colophon

- Some slide material is from Stanford CS145 (Jeffrey D. Ullman, Fall 2007)