

3. Updates

- To change specific attributes in certain tuples of a relation:

```
update <relation>  
set <list of attribute assignments>  
where <condition on tuples>;
```

- Example: Change patron Nigel's phone number to 250-214-9999

```
update Patrons  
set phone = '250-214-9999'  
where name = 'Nigel';
```

Update several tuples

- Make \$4 the maximum price for beer:

```
update sells  
set price = 4.00  
where price > 4.00;
```

Constraints & Triggers

- **Constraint**
 - A relationship among data elements the DBMS is required to enforce
 - Example: key constraints
- **Triggers**
 - A technique for implementing some types of constraints
 - Only executed when a certain condition or event occurs
 - Example: Trigger executed when a tuple is inserted into a table.
 - Rationale: Triggers are a way to implement seemingly complex constraints.
- We will look at each of these mechanisms in some detail

Constraints: different kinds

- **Keys**
- **Foreign keys**
 - Also referred to as referential integrity constraints
- **Value-based** constraints
 - Constrain values for a particular table attribute
- **Tuple-based** constraints
 - Apply to relationships amongst components in the database
- **Assertions**
 - Anything that can be denoted as an SQL boolean expression

Recall: single-attribute keys

- Used the keywords **primary key** or **unique**
- These are placed after the statement of type in a table declaration

```
create table Beers (  
    name    char(20)    unique,  
    manf    char(20)  
);
```

Recall: multi-attribute keys

- The **bar** and **beer** attributes together are the key for the Sells relationship

```
create table Sells (  
    pub      char(20),  
    beer     varchar(20),  
    price    real,  
    primary key (pub, beer)  
);
```

Foreign keys

- Idea:
 - Values appearing in attributes of one relation **must appear** together as specific in another relation.
- Example:
 - Consider **Sells(pub, beer, price)**
 - We might expect that a **beer** value also appears in **Beers** as **name**
 - (Note: Here we are ourselves asserting the degree to which our model must correspond to the real-world equivalent.)

Foreign keys: expressing

- Use the keyword **references**
 - After an attribute (for one-attribute keys), or
 - As an element of the schema.
- Attributes referenced in the other schema must themselves be declared primary key or unique

```
foreign key (<list of attributes>)  
references <relation> (<attributes>)
```


Example: denoted with attribute

```
create table Beers (  
    name    char(20) primary key,  
    manf    char(20)  
);
```

```
create table Sells (  
    pub      char(20),  
    beer     char(20) references Beers(name),  
    price    real  
);
```

Example: denoted as schema element

```
create table Beers (  
    name    char(20) primary key,  
    manf    char(20)  
);
```

```
create table Sells (  
    pub      char(20),  
    beer     char(20),  
    price    real,  
    foreign key(beer) references Beers(name)  
);
```

Foreign keys: enforcing the constraints

- If there is a foreign-key constraint from relation **R** to relation **S**...
- ... then **two violations are possible**.
 1. An **insert** or **update** to **R** introduces values not found in **S**.
 2. A **deletion** or **update** to **S** causes some tuples of **R** to "dangle"

Foreign keys: actions taken

- Example:
 - Suppose R corresponds to Sells ...
 - ... and S corresponds to Beers
- Possibility 1: An insert or update to Sells that introduces a non-existent beer must be rejected.
- Possibility 2: A deletion or update to Beers removing a beer value found in some tuples of Sells can be handled by:
 - **Rejecting** ...
 - **Cascading** ...
 - ... or **Nulling**.

Foreign keys: actions taken

- Default: **rejecting**
 - modification is not applied (i.e., it is rejected)
- **Cascading**: Same change in Beers is made in Sells
 - Updated beer: change all corresponding values in Sells
 - Deleted beer: delete all corresponding Sells tuples
- **Null**: Change beer in Sells tuples to the **null** value.

Cascade: examples

- Recall: $R == \text{Sells}$, $S == \text{Beers}$
- Delete from S
 - Delete the ('Blue', 'Labatt's') tuple from Beers.
 - Result: DBMS delete all tuples from Sells where beer = 'Blue'
- Update to S
 - Change the ('Blue', 'Labatt's') tuple in Beers to ('Blue (Pilsner)', 'Labatt's').
 - DBMS changes all Sells tuples with beer = 'Blue' to 'Blue (Pilsner)'

Set null: examples

- Delete ('Blue', 'Labatt's') tuple from Beers
 - DBMS changes all tuples of Sells with beer = 'Blue'...
 - ... to beer = null.
- Update the ('Blue', 'Labatt's') tuple in Beers
 - Same change as in previous slide...
 - ... results now results in affects Sells tuples having null for the beer attribute.

Choosing a policy

- To choose reject, cascade, or update...
- ... we do so at table creation time ...
- ... and also select specific policy for action (i.e., update vs. delete).
- Follow the foreign-key declaration with **on [update, delete] [set null cascade]**
 - Two such clauses may be used
- With no such clause, default (i.e., reject) is used.

Example: specifying policy

```
create table Sells (  
    pub      char(20),  
    beer     char(20),  
    price    real,  
    foreign key(beer)  
        references Beers(name)  
        on delete set null  
        on update cascade  
);
```