

Trigger: the condition

- The condition appears within a function
 - Also known as a **stored procedure**
- This condition is normally boolean valued
 - However, as the stored procedure can have arbitrary code, it might not check any condition.
 - For our purposes, however, we do want to check it.
- Condition is evaluated on the database as it would exist:
 - Before the triggering event (if **before** is specified in the **create trigger** statement)
 - After the triggering event (if **after** is specified in the **create trigger** statement)
 - Regardless, applied before the changes take effect!
 - Access tuple via NEW and OLD variables (if using **for each row**)

Trigger: the action

- An action can consist of more than one SQL statement
 - These are normally within the true branch of the procedure's if statement
- Note:
 - Queries make no sense in an action...
 - ... so we are really concentrating on modifications.

Trigger: second example

- Consider Sells(pub, beer, price)...
- ... and NastyPubs(pub)
- We want to maintain a list of pubs that raise the price of any beer by more than \$1
- Idea:
 - We examine when a row in Sells is updated...
 - ... such that its old price differs from its new price ...
 - ... where the new price is $> \text{old price} + 1.00$.
 - If that happens, we insert the pub into NastyPubs

Trigger: condition + action; create trigger

```
create function price_trig_func() returns trigger as $price_trig$
begin
    if NEW.price > (OLD.price + 1.00) then
        insert into NastyPubs(pub) values (NEW.pub);
    end if;
    return NEW;
end;
$price_trig$ language plpgsql;
```

Assuming here that the table NastyPubs has already been created.

```
create trigger PriceTrig after update on Sells
for each row
execute procedure price_trig_func();
```

Trigger: third example

- Some triggers are not applied to specific rows
 - Rather, they apply to a given SQL modification statement that might affect one or many rows
 - We use **for each statement** to indicate that the trigger is performed not for each modified row but once a whole statement is completed
- Example: Keeping track of stats
 - If there is a change to the beers sold in pubs (i.e., in Sells), then update PubStats to reflect the changed average price of all beers sold in all pubs

We will need this...

```
create table PubStats (  
    scope          varchar(20),  
    avgbeerprice  real,  
    numberpubs    int,  
    newestpub      varchar(20)  
);  
insert into PubStats(scope) values ('all');
```

Trigger: condition + action; create trigger

```
create function average_trig_func() returns trigger as
$average_trig$
begin
    update PubStats
        set avgbeerprice = (select avg(price) from sells)
        where scope = 'all';
    return null;
end;
$average_trig$ language plpgsql;
```

```
create trigger UpdateAverageTrigger
after update or delete or insert on Sells
for each statement
execute procedure average_trig_func();
```

A bit of system-design philosophy...

- We have seen some constraints which are attribute- or tuple-based
- We have seen others which link a stored procedure to some events (ECA rules)
- Skipping ahead:
 - We also know some of these constraints could, in some way, be enforced by our own code
 - That is, the logic of our implemented algorithms would ensure data consistency
- Where do we state/implement a constraint?
 - In the database? I
 - In our code?
 - A mixture of both?

Transactions

- Up to this point:
 - We have concentrated on a single database
 - We have assumed a single user performs all queries and modifications
 - Also have assumed that user does these sequentially (i.e., one after the other)
- However, these assumptions can be too strict
 - Database systems are normally accessed by many users or processors (or both) at the same time
 - These accesses mix both queries and modifications
- While operating systems support process interaction...
- ... DBMSes need some help to control potentially trouble accesses as they interact.

Troublesome example

- Your name is Pat
- Your partner's name is Chris
- You each transfer \$500 from your joint account to your own individual accounts...
 - ... and do so from different ATMs...
 - ... at the same time ...
 - ... and not because you are arguing with each other.
- What we expect:
 - DBMS supporting the bank's operations does not lose one of the transactions.
- What we don't expect:
 - Classic race condition a la operating systems
 - Lost / propagating money
 - Need for couple's counselling

Potential order of operations

```
select balance from SavingsAccounts  
where customer = 'Pat and Chris';
```

```
-- Bank's computer checks if balance is > 500; if so, then...
```

```
update SavingsAccount set balance = balance - 500  
where customer = 'Pat and Chris';
```

```
update ChequingsAccount set balance = balance + 500  
where customer = 'Pat';
```

```
select balance from SavingsAccounts  
where customer = 'Pat and Chris';
```

```
-- Bank's computer checks if balance is > 500; if so, then...
```

```
update SavingsAccount set balance = balance - 500  
where customer = 'Pat and Chris';
```

```
update ChequingsAccount set balance = balance + 500  
where customer = 'Chris';
```

Inappropriate order of operations

```
select balance from SavingsAccounts  
where customer = 'Pat and Chris';
```

```
-- Bank's computer checks if balance is > 500; if so, then...
```

```
select balance from SavingsAccounts  
where customer = 'Pat and Chris';
```

```
-- Bank's computer checks if balance is > 500; if so, then...
```

```
update SavingsAccount set balance = balance - 500  
where customer = 'Pat and Chris';
```

```
update ChequingsAccount set balance = balance + 500  
where customer = 'Pat';
```

```
update SavingsAccount set balance = balance - 500  
where customer = 'Pat and Chris';
```

```
update ChequingsAccount set balance = balance + 500  
where customer = 'Chris';
```

Transactions

- There are many variants of this for different situations
 - State of DBMS indicates a possible change can be made
 - Change is performed
 - Want system in some "consistent" state (from the point of view of the users) after the change
- **Transaction**
 - A process involving database queries or modification or both
 - Normally with strong properties regarding concurrency
 - Formed in SQL from single statements or explicit programmer control.

Properties: ACID

- Transactions are normally designed to be **ACID**
- **Atomic**: Either the whole transaction is performed, or not if it is.
- **Consistent**: Database constraints are preserved.
- **Isolated**: It appears to the users as if only one processes executes at any one time.
- **Durable**: Effects of a process will survive a crash of the system
- Optional:
 - Weaker forms of transaction are also supported (**serializable level**)
 - These permit a higher degree of concurrency but with some potential problems.

Transaction: syntax

- Sequence of database actions are bracketed by **begin** and **commit**
- After the **commit** operation, database modifications are permanent in the database

```
begin;  
  
update accounts set bal = bal - 90000 where name = 'Nigel Wright';  
update accounts set bal = bal + 90000 where name = 'Michael Duffy';  
update accounts set bal = bal - 90000 where name = 'Michael Duffy';  
update accounts set bal = bal + 90000 where name = 'Senate Petty Cash';  
  
commit;
```

Rollback

- A transaction can be ended with **rollback** instead of **commit**...
- ... but this instead causes the transaction to be aborted rather than results written to the DBMS
- That is:
 - rollback ensure there are no changes to the database caused by the aborted transaction
- Note:
 - Some failures can require a rollback even without the programmer knowing it.
 - Example: division by 0
 - Example: constraint violation

Interacting processes: example

- Consider our relation Sells(pub, beer, price)
- Suppose the Valhalla Inn pub sells **only** the following two beers: Blue for \$4.00 and Coors Light for \$4.25.
- Siegfried is querying Sells for the highest and lowest prices charged by Valhalla Inn
- The manager at Valhalla Inn decides to stop selling Blue and Coors Light...
- ... but she adds Corona to the menu for \$4.50.

Program: run by Siegfried

- Siegfried performs the following two SQL statements.
- The result from each query is stored some place by Siegfried

```
-- max operation  
select max(price) from Sells  
where pub = 'Valhalla Inn';
```

```
-- min operation  
select min(price) from Sells  
where pub = 'Valhalla Inn';
```

Program: run by Brunnhilde from the pub

- Brunnhilde (manager at the Valhalla Inn) executes the following delete and insert steps.
- Why she has write-access to the table used by all pubs is a technical detail we will gloss over for now.

```
-- del operation  
delete from Sells  
where bar = 'Valhalla Inn';
```

```
-- ins operation  
insert into Sells  
values ('Valhalla Inn', 'Corona', 4.50);
```

Interleaving of statements

- We know that:
 - the **max operation** occurs before the **min operation**
 - the **del operation** occurs before the **ins operation**
- There are at present no other restrictions on the order of the operations
- In order to add restrictions to Siegfried and Brunnhilde's actions...
 - ... We would need to use transactions.
 - (Restrictions here only apply to the actions related to the database.)

Example: strange interleaving

- Assume we are not yet using transactions.
- Suppose the operations are executed in this order:
 - (max)(del)(ins)(min)
- (max): Set of prices is {4.00, 4.25}, the max is 4.25
- (del): Set of prices becomes {}
- (ins): Set of prices becomes {4.50}
- (min): As set of prices is {4.50}, the min is 4.50
- Siegfried sees that the $\text{max} < \text{min}$!
 - This is not good (i.e., this is an example of **inconsistency**)