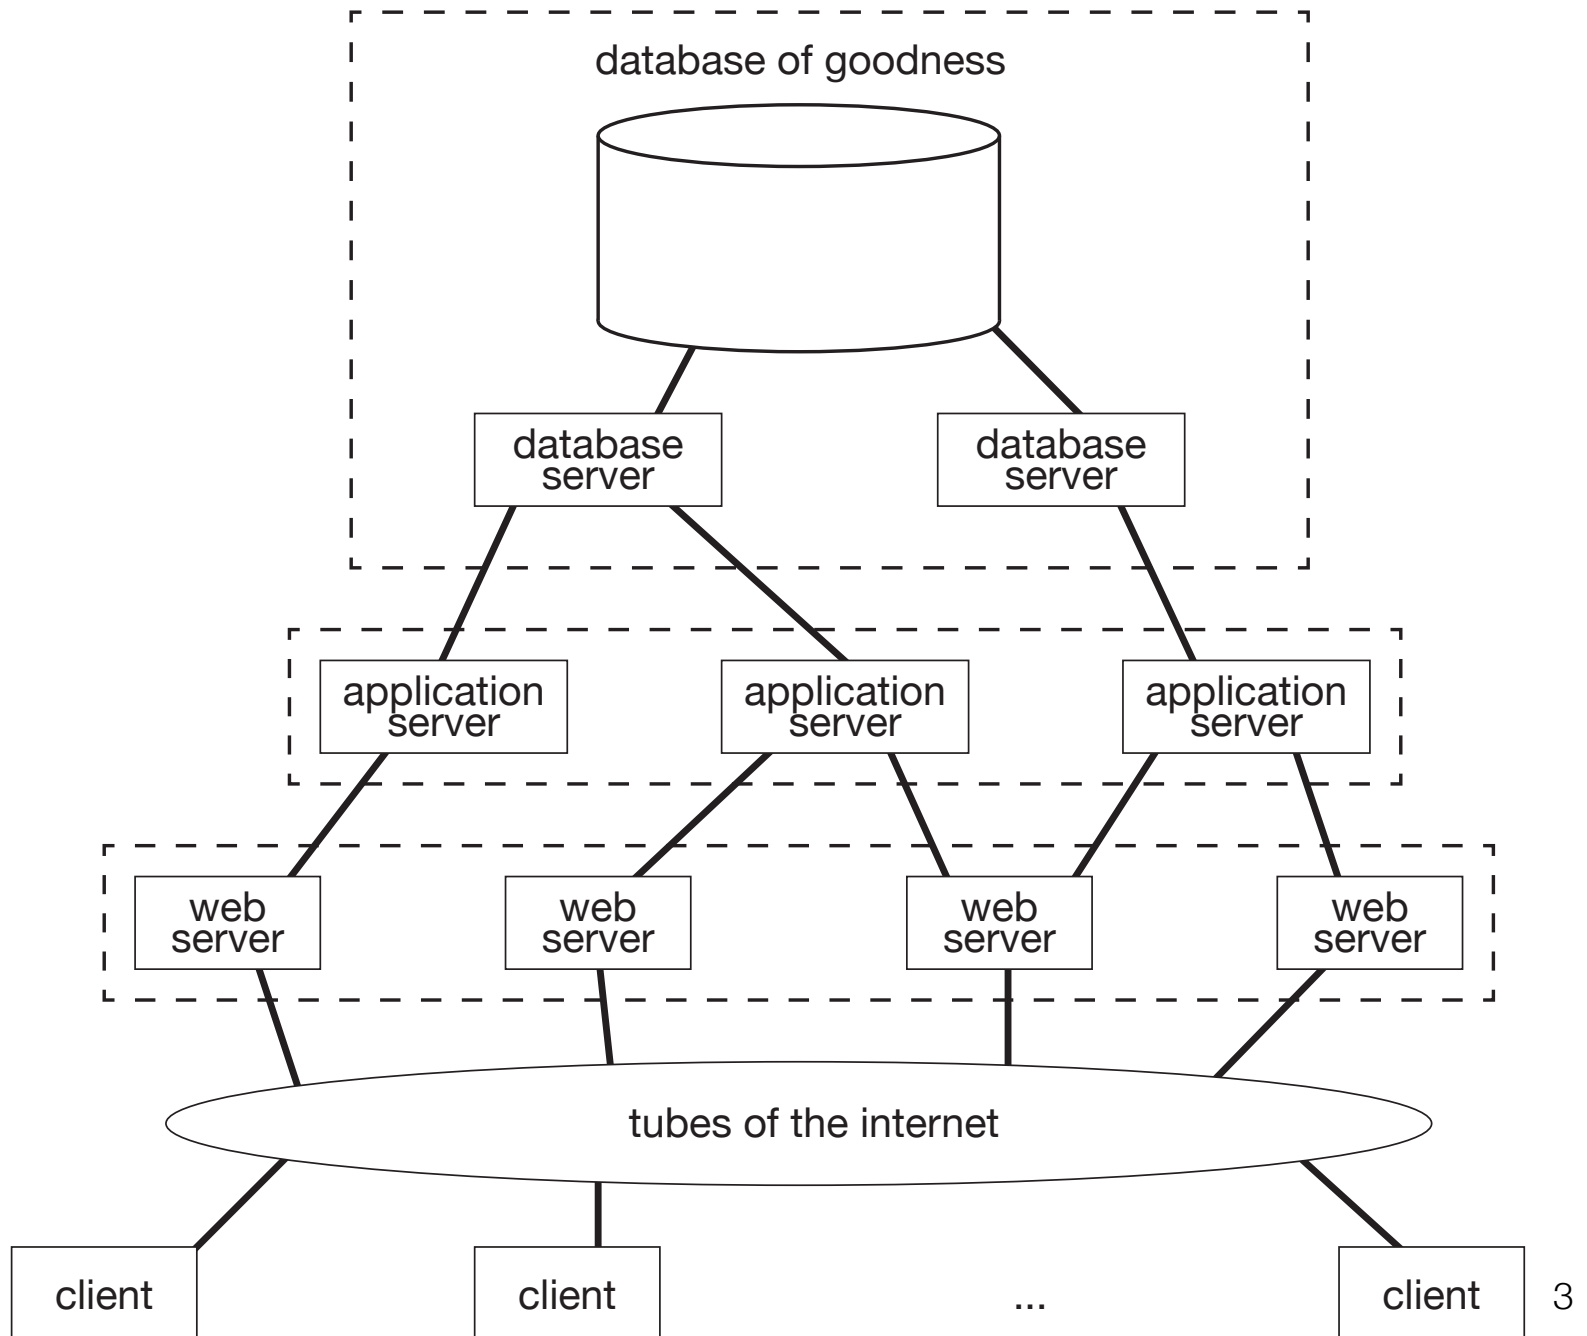


CSC 370

Programmatic access to a DBMS

Present coding reality

- Many production systems are multi-tiered
 - Front-end: web-server tier
 - Middle-tier: application tier
 - Bottom-tier: database tier
- At this point in the course we are interested in how to access a DBMS programmatically
 - The result is not necessarily sexy...
 - ... as it often involves a mix of SQL and regular code.
 - However, what we will use will hopefully look better than PHP!



Web-server tier

- Manages interactions with the user
- Often running Apache/Tomcat
 - Helps handle details such as session management
 - Users connecting to the webserver via their browser are **clients**
- Example: www.amazon.ca
 - User opens a connection to the Amazon database by first entering the URL of the site
 - Webserver presents a view of the system (home page) which is rendered by the user's browser
 - Forms, menus, buttons enable client to express what they want accomplished
 - Client's browser transmits the information to the webserver
 - Webserver must then negotiate with the next tier (application-server tier)

Application-server tier

- Turns data (from database) into a response to a webserver request (ultimately from the client)
- A web-server process can invoke one or more application-tier processes to handle the request
 - Actions performed by application tier: sometimes called business logic
 - In essence, application server is designed to respond to a customer's request and implement the company's strategy for handling that request
- Amazon.ca again: Possible response are...
 - Show the user's wishlist
 - Fetch information about a book
 - Add book to a cart
- May also contain a subtier ("database integration")

Database tier

- Executes queries requested from the application-server tier
 - May also buffer data needed to handle some requests
- In a large system, database tier may consist of many servers
 - Query load is distributed across the servers
 - System must ensure association between application server and database-server instance is maintained
- Connections must be established to a database server...
 - ... and these are expensive (in time) to establish.
 - Therefore some management of connections (pools) is often used.

Our interest at this point

- We are concerned now with how to implement the database tier
 - Implementing the webserver and application tiers is a topic for a course such as SENG 450.
 - Note: some programming frameworks help combine work on all three tiers into one package (e.g., Django, Ruby on Rails)
- For now, we need to learn:
 1. How does a database interact with regular programs?
 2. How can we deal with data-type differences between SQL and our chosen language?
 3. How do we manage connections to short-lived processes?

psycopg2 (note spelling – no "h"!)

- Database adapter for Python
 - There are several, but this one is well use.
 - Is also available on the lab machines.
 - Chat with instructor if you need more information on how to install it on your own machines.
- Literally a wrapper...
 - ... as it is a way of accessing the official libpq Postgres client library.
- We will use it to access your database on studentdb.csc.uvic.ca
- Can of course be used to access other PostgreSQL servers
 - But remember that any connections to studentdb.csc.uvic.ca must originate from within the Faculty of Engineering's network.

psycopg2: Our "hello, world!"

```
#!/usr/bin/python

import psycopg2

def main():
    dbconn = psycopg2.connect(host='studentdb.csc.uvic.ca', user='c370_s63',
                              password='fS8qPPLc')
    cursor = dbconn.cursor()

    cursor.execute("""
select *
from sells
where price < 5
""")

    for row in cursor.fetchall():
        print "%s %s %s" % (row[0], row[1], row[2])

    cursor.close()
    dbconn.close()

if __name__ == "__main__": main()
```

psycopg2: Our "hello, world!" output

```
$ ./hello_world.py
Bard & Banker      Blue      3.5
The Hacked Library Blue      3.25
The Hacked Library Bud Light 4.0
```

Python and database APIs

- Python has a mixed reputation with database folks
 - Language is great to use...
 - ... but adapter support has not been consistent.
- However:
 - There is a standard Python database API
 - (PEP 249 – now at version 2.0)
 - Standard describes what is needed in an adapter.
 - Must also be usable by all major DBMS implementation (i.e., might not be ideally tuned for Postgres)
- We will be looking at some of the simpler functionality
 - And don't forget – Postgres permits us to create stored procedures, so not everything need be done via Python.

A word about driver independence

- Often best to write code that is not specific to a DBMS or driver
 - Makes it easier when porting the program to a different platform
 - Makes it easier when porting program to use a different DBMS
- Relatively straightforward to do this in Python

Accessing driver outside of "import"

```
#!/usr/bin/python

DBTYPE = 'psycopg2'

def main():
    dbdriver = __import__(DBTYPE, globals(), locals(), [], -1)

    dbconn = dbdriver.connect(...)

# etc. etc. etc.
```

Setting up a connection

- Everything that holds for connecting to the studentdb server holds for connecting via Python
 - Must be connecting from a machine in the Faculty of Engineering network
 - Provide username and password
- There are 2.5 different ways to set up the connection object
 - Using libpq connection string
 - Using keyword arguments

Connection approaches

libpq-ish string

```
dbconn = psycopg2.connect("host=studentdb.csc.uvic.ca user=c370_s63 "  
    + "password='fS8qPPLc'")
```

libpq-ish with explicitly named argument

```
dbconn = psycopg2.connect(dsn="host=studentdb.csc.uvic.ca user=c370_s63 "  
    + "password='fS8qPPLc'")
```

keyword arguments

(but not everything libpq is necessarily supported by psycopg2)

```
dbconn = psycopg2.connect(host='studentdb.csc.uvic.ca',  
    user='c370_s63',  
    password='fS8qPPLc')
```

For more information on libpq (the C interface to Postgres) go to

<http://www.postgresql.org/docs/9.1/static/libpq.html>

Cursors

- Some mechanism is needed to permit programmatic access to the results of a query
 - The query result is a relation...
 - ... which is made up of rows with attributes ...
 - ... although there is no guarantee these rows easily map to a program data structure
- To provide consistency, database-interface languages support the concept of a **cursor**
 - A control structure that enables the code to traverse the rows in the relation.
 - In essence it supports operations such as retrieval, addition and deletion of tuples / table rows.
 - (Acts a bit like an iterator)
 - (How much data should be transferred?)

Cursor: style 1

- A cursor is associated with a connection for the life of that connection
- If unnamed, the cursor object really isn't a cursor object
 - Rather, it is handle to data stored within the Python program (a "statement handle")
 - That is, all the data is transferred from the server to the client!

```
cursor = dbconn.cursor()
```

Cursor: style 2

- **server-side cursor**
 - Controls that amount of data transferred from server to client.
 - Very useful when there is a large amount of data.
 - For this to work, our code will need to fetch data from the cursor from time to time.
- Also called a named cursor
- (Semantics can be a bit confusing when cursors and transactions mix...)

```
cursor = dbconn.cursor(name="mycursor")
```

Executing

- Cursors permit:
 - querying
 - update modifications
 - other Postgresql operations

```
cursor.execute("""  
select pub, beer  
from sells  
where price < 5  
""")
```

```
cursor.execute("update sells set price = 3.00 where beer = 'Blue'")
```

```
cursor.execute("analyze verbose sells")
```

Fetching results

- Traversing through the cursor results is known as fetching
 - fetchall
 - fetchmany
 - fetchone
- Beware: not everything can be fetched!

```
cursor.execute("""
select pub, beer
from sells
where price < 5
""")
result = cursor.fetchall()
print result

cursor.execute("analyze verbose sells")
result = cursor.fetchall()
print result
```

fetchall output: examples

```
$ ./slide15.py
[('Bard & Banker', 'Blue'), ('The Hacked Library', 'Blue'), ('The Hacked Library', 'Bud Light')]

Traceback (most recent call last):
  File "./slide15.py", line 26, in <module>
    if __name__ == "__main__": main()
  File "./slide15.py", line 23, in main
    result = cursor.fetchall()
psycopg2.ProgrammingError: no results to fetch
```

Fetching results

- Iterating through each result row is straightforward...
 - ... but we do lose some information.
 - Attributes must be accessed by position in the row (i.e., not by their name).

```
cursor.execute("select pub, price, beer from sells")
for row in cursor.fetchall():
    print "Sells data: %s %s %s" % (row[0], row[1], row[2])
```

```
$ ./slide17.py
Sells data: Bard & Banker      3.5 Blue
Sells data: The Hacked Library 3.25 Blue
Sells data: The Hacked Library 7.0 Amnesiac
Sells data: The Hacked Library 4.0 Bud Light
```

More fetching techniques

- Fetch a single row
 - Make sure, however, that there was a row returned
 - (Unlike an iterator loop, need to use an "if")

```
cursor.execute("select pub, price, beer from sells")
row = cursor.fetchone()
if row is not None:
    print "Sells data: %s %s %s" % (row[0], row[1], row[2])
```

Fetching in batches

- If there is a lot of data at the server, we may only want to process a bit of it at a time
- We can control the maximum number of rows returned
 - Specify the maximum number of batch rows before we execute the SQL statement

```
cursor = dbconn.cursor()
cursor.arraysize = 500
cursor.execute("select shrt_desc, ndb_no from food_des order by shrt_desc")

batch_num = 1
while True:
    batch = cursor.fetchmany()
    if not batch:
        break
    print "BATCH # %d" % (batch_num)
    for row in batch:
        print "Desc: %s; Database #: %s" % (row[0], row[1])

    batch_num = batch_num + 1
```


fetchmany example: output

```
$ ./slide19.py
BATCH # 1
Desc: ABALONE,MIXED SPECIES,RAW; Database #: 15155
Desc: ABALONE,MXD SP,CKD,FRIED; Database #: 15156
Desc: ABIYUCH,RAW; Database #: 09427
Desc: ACEROLA JUICE,RAW; Database #: 09002
Desc: ACEROLA,(WEST INDIAN CHERRY),RAW; Database #: 09001

<... stuff deleted ...>

Desc: WALRUS,LIVER,RAW (ALASKA NATIVE); Database #: 35083
Desc: WALRUS,MEAT,DRY (ALASKA NATIVE); Database #: 35079
BATCH # 15
Desc: WALRUS,MEAT,RAW (ALASKA NATIVE); Database #: 35081
Desc: WALRUS,MEAT & SUBCUTANEOUS FAT (ALASKA NATIVE); Database #: 35082

<... stuff deleted ...>

Desc: YOGURT,VAN OR LEM FLAV,NONFAT MILK,SWTND W/LOW-CALORIE SWTNR; Database #:
01184
Desc: ZWIEBACK; Database #: 03217
$
```

Metadata

- Even though we must access attributes by row index...
 - ... we can obtain some table metadata from the cursor.
- Can use this metadata as appropriate.
 - Not a lot we can do with type_code besides compare for equality.

```
cursor.execute("select price, pub, beer from sells")
print cursor.description[0].name
print cursor.description[0].type_code
print cursor.description[2].name
print cursor.description[2].type_code
```

```
$ ./slide21.py
price
700
beer
1042
```