

B+ tree: wrapping up for now


- B+ deletion works by merging nodes to avoid underflow
 - Deletion is also recursive.
 - In practice, however, we can simply delete the key from the leaf...
 - ... as we are often accurate in assuming that databases grow more often than they shrink.
- B+ tree can be used as a sorting algorithm for disk-based sorting.

Hash Tables

- As applied to indexes, these are variants of what you have already seen in other courses:
 - Hash function $h()$: takes a search key and computes an integer in the range of 0 to $B-1$
 - Bucket array: array of B linked lists.
 - A $\langle \text{key}, \text{value} \rangle$ pair is to be stored in bucket numbered $h(\text{key})$
- These hash tables must reside on disk (i.e., too large for main memory)
 - Each bucket is, therefore, the size of a disk block
 - Buckets can be extended via overflow blocks
- With a good hash function and number of buckets, we expect behavior comparable to a B+ tree.


Hash table: example

- Buckets are numbered from 0 to 3
- Keys are letters
- Two <key, value> pairs per bucket.
 - Hash function here is indicated (e.g., $h(a)=3$, $h(b)=2$, $h(c)=1$, etc.)
- Assume some in-memory structure maps from hash number to disk block (bucket)
- Additional information for each block is represented by a small tab.

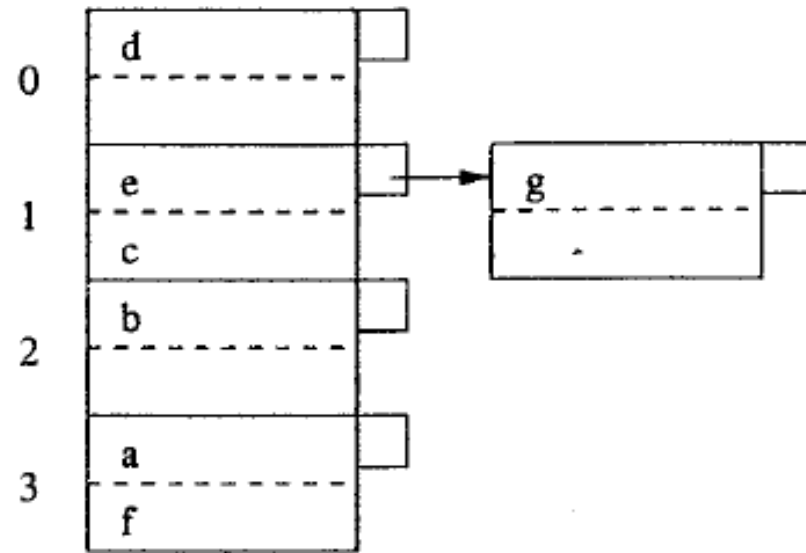


Hash table: insert key/value with key "g"

$$h("g") = 1$$




before




after

Hash table: delete key/value with key "c"



before



after

Static vs. dynamic hash tables


- Where the number of buckets does not change:
 - We have a **static hash table**
- A problem with hash tables is the degeneration of buckets into long linked lists
 - Traversing through the linked list would be one I/O per overflow block
- Recall: want to keep I/O cost low
 - Hence want to keep the block-per-bucket ratio low.
 - Best way to do that: keep an optimal number of buckets at all time
- **Dynamic hash table:**
 - B is allowed to vary so that there is about one block per bucket.

Extensible hash table

- (Sometimes also called **extendible hashing**.)
- A kind of dynamic hash table
 - Indirection (via pointers) used to access buckets
 - Array of pointers to buckets can grow.
 - There need not be a disk block per bucket (i.e., certain buckets can share a block)
 - Hash function **$h()$** computes a sequence of **k** bits for the key
 - Some buckets use a smaller number of the k bits, some a larger number (i.e., a number **i** for each bucket is maintained where $i \leq k$)

Small extensible hash table

- Hash function produces a sequence of four bits
 - $k = 4$
- Only one of the bits is used at present
 - Indicated by the box above the bucket pointers where $i = 1$
 - We use the most significant bits here.
- Relationship between i and k :
 - First bucket holds all keys whose hash value begins with 0
 - Second bucket holds all keys whose hash value begins with 1.



Extensible hash table insertion

- Goal:
 - Insert $\langle K, V \rangle$ pair into an appropriate bucket...
 - ... yet make sure we add buckets only when necessary ...
 - ... and adjust the number of key bits used to map to the bucket.
- Recall:
 - $h(K)$ is a bit sequence
 - i indicates how many of the first bits of $h(K)$ are used to locate a bucket
 - i is stored with the bucket
- Three cases:
 - A. There is room in the bucket for $\langle K, V \rangle$.
 - B. There is no room in the bucket, and to locate a key-value pair we need fewer bits than i .
 - C. There is no room in the bucket, and to locate a key-value pair we need exactly i bits.

Case B

```
k = h(K)
B = bucket_array.forkey(k)
j = identifying_bits(B, k)
```

- There is no room in the bucket, and $j < i$ (where j is stored in that "tab" attached to a bucket).


```
B0 = new Bucket()
B1 = new Bucket()

for each item in B:          # Move items from B into either B0 or B1
    if item.key.bit(j+1) == 0:
        B0.insert(item)
        bucket_array.setkey(B0, item.key())    # Bucket array now points B0
    else:
        B1.insert(item)
        bucket_array.setkey(B1, item.key())    # Bucket array now points to B1
```


Case C

- There is no room in the bucket, and to locate a key-value pair we need exactly i bits.
- In all of our examples so far, we haven't shown how this global value of i is changed.
 - A change was sneaked in
- With this case, we must increment i , possibly several times
 - Incrementing i means doubling the number of bucket pointers (from 2^i to 2^{i+1})
 - Must update pointers to buckets as some pointer pairs still lead to the same bucket.
- We only split a bucket when we have enough bucket pointers to that case B takes over.

Example: Insert $\langle K, V \rangle$ where $h(K) = 1010$




before




after

Insert pairs with keys 0000, 0111, 1000



before



after