

Linear hash table

- This is another dynamic hash table approach
 - Bucket number grows more slowly than with extensible hashing
 - Fill factor: Number of buckets **n** always chosen so average number of <key, value> pairs (value **r**) is some constant fraction of **n** (e.g. $r \leq 1.7n$)
 - Blocks cannot be split, so overflow blocks are used instead (but average # of overflow blocks will be much less than 1)
 - Number of bits used number bucket entries is $\text{ceiling}(\log_2 n)$

Small linear hash table

- Note the three global values
 - i : number of bits used to determine bucket for key
 - n : number of buckets
 - r : number of entries in the hash table
- Unlike extensible hashing, we use least-significant bits of key to compute bucket location.


$i=1$
$n=2$
$r=3$

0	0000	
	1010	
1	1111	


Linear hash table insertion

- Goal:
 - Compute the bucket in which the $\langle \text{key}, \text{value} \rangle$ pair should be placed.
 - If there is room in the bucket, great!
 - If no room, then create an overflow bucket.
 - If occupancy exceeds the fill factor, then create a new bucket.
- Computing bucket looks more complicated than it is.
 - Given some key K ...
 - ... denote its least-significant i bits as $a_1 a_2 \dots a_i$
 - ... and call this bit sequence **m** .
 - If $m < n$: bucket numbered m exists, and place key-value pair in that bucket.
 - If $n \leq m \leq 2^i$ then bucket m does not yet exist, so put key-value pair into bucket $m - 2^{i-1}$ (i.e., same as setting a_1 to 0)

Insert pair with key 0101




before

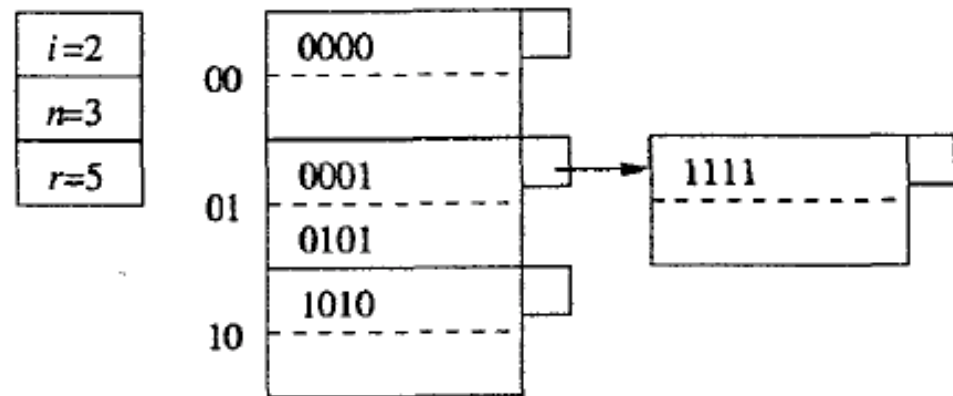


after

Insert pair with key 0001

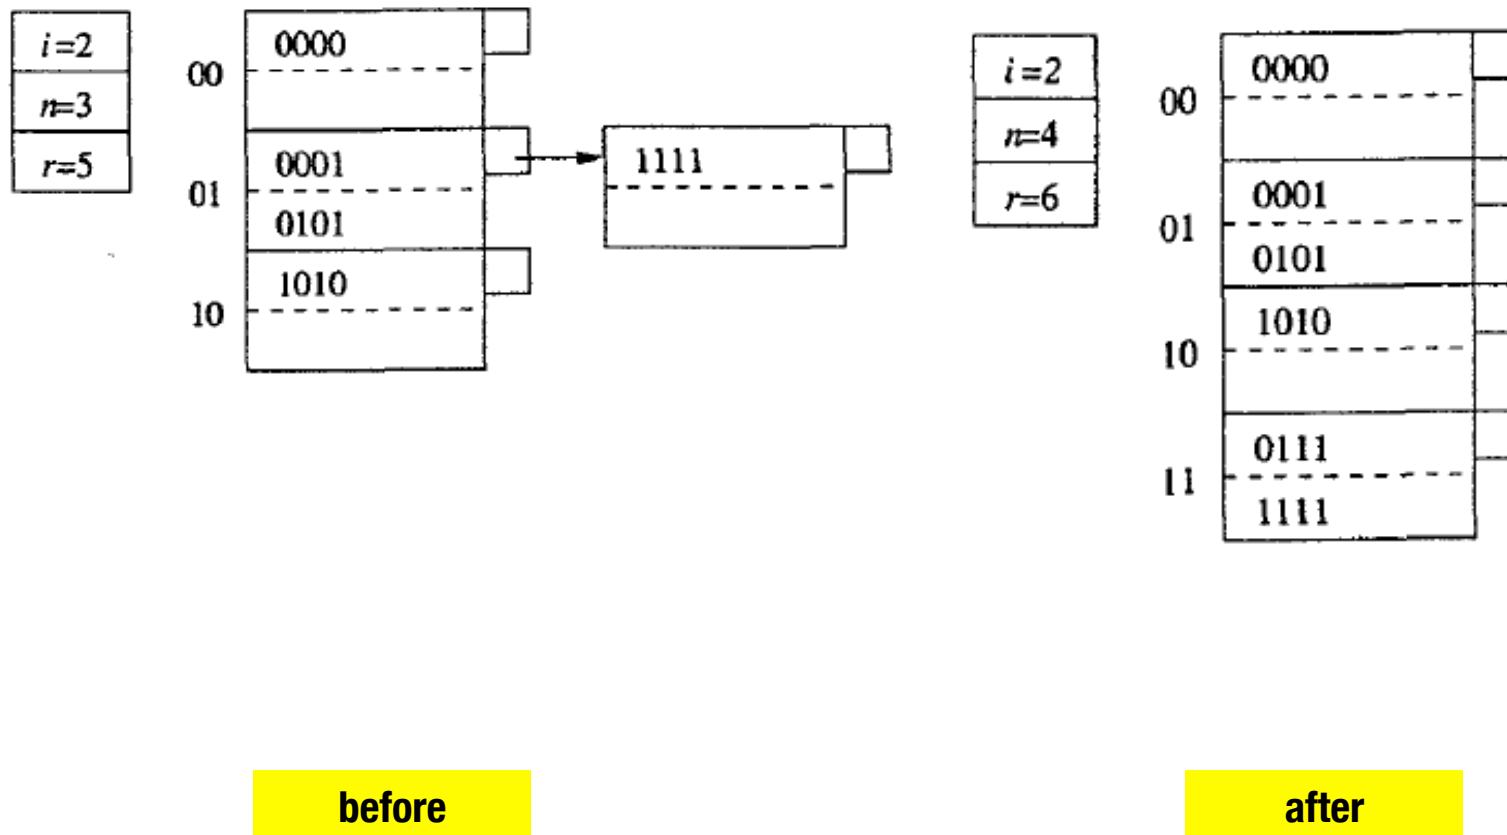


before



after

Insert pair with key 0111



Story so far

- All of the indexes so far have been one-dimensional
 - Each has a single search key (which may comprise one or more table attributes)
 - Values for all attributes of the search key must be provided.
 - Index search is through this sequence of values for a matching index key (and corresponding data-file block).
- B+ Trees:
 - Single linear ordering for keys
- Hash tables:
 - Search key is completely known for lookup


Multidimensionality

- Many current applications rely on the use of **multidimensional data**
 - Geosciences; cartography
 - Mechanical CAD; VLSI CAD; 3D CG modelling
 - Robotics
 - Visual perception
 - Autonomous navigation
 - Environmental protection
 - Medical imaging
- May also refer to this as **spatial data**

What is so special about multidimensional data?


- Complex structure
- Often dynamic
- Tend to be large
- No standard algebra on multidimensional data
- Operators are not closed
- Computational costs vary among multidimensional database operators

Kinds of multidimensional queries (1)



Point query


$$PQ(p) = \{o \mid p \cap o.G = p\}$$



Window query (also Range query)


$$WQ(I^d) = \{o \mid I^d \cap o.G \neq \emptyset\}$$

Kinds of multidimensional queries (2)



Intersection Query


$$IQ(o') = \{o \mid o'.G \cap o.G \neq \emptyset\}$$



Enclosure query


$$EQ(o') = \{o \mid (o'.G \cap o.G) = o'.G\}$$

Kinds of multidimensional queries (3)



Containment Query

$$CQ(o') = \{o \mid (o'.G \cap o.G) = o.G\}$$



Adjacency Query

$$AQ(o') = \{o \mid o.G \cap o'.G \neq \emptyset \\ \wedge o'.G^\circ \cap o.G^\circ = \emptyset\}$$

Other queries

- Exact match query
 - $\text{EMQ}(o') = \{o \mid o'.G = o.G\}$
- Nearest-neighbour query
 - $\text{NNQ}(o') = \{o \mid \forall o'' : \text{dist}(o'.G, o.G) \leq \text{dist}(o'.G, o''.G)\}$
- Spatial join
 - R and S are collections of multidimensional objects
 - $R \bowtie_{\theta} S = \{(o, o') \mid o \in R \wedge$
 $o' \in S \wedge$
 $\theta(o.G, o'.G)\}$
 - θ is a spatial predicate (e.g., intersects, contains, is_enclosed_by, distance, etc.)